

Time Series Datasets for Anomaly Detection in R



Data Analytics Lab

dal.global@grupo.cefet-rj.br | <https://eic.cefet-rj.br/dal>



Outline

This collection presents the benchmarks distributed through Harbinger's unified integration layer. The goal is practical: understand what each benchmark contains **before** using it in anomaly detection, change-point, or motif discovery experiments.

01

Benchmark Files

Yahoo S5, NAB, and UCR — widely used reference collections for time series analysis

02

Domain Collections

MIT-BIH, 3W Oil, and GECCO — real-world signals with rich applied context

03

Recommended Workflow

First understand the collection structure, then select the appropriate family of techniques

📄 Run the full examples and explore the source code at: <https://github.com/cefet-rj-dal/harbinger/tree/master/examples>

Inspect Before You Model

Before applying any event detection technique to a time series dataset, it's essential to understand its underlying structure. This "inspect first" discipline prevents mismatches between the dataset's characteristics and the chosen detection method.

Dataset Structure

In Harbinger, a dataset is a list of data frames, each representing one time series.

Column Separation

Metadata columns (`idx`, `event`, `type`, `seq`, `seqlen`) are separated from the actual signal columns.

Data Type

Classify as univariate or multivariate based on the number of signal columns present in the data.

Visualize First

Always perform a visual inspection (plot preview) of your time series data before proceeding with modeling.

`dataset_summary()`

```
dataset_summary <- function(x) {  
  first_series <- x[[1]]  
  meta_cols <- c("idx", "event",  
    "type", "seq", "seqlen")  
  signal_cols <- setdiff(  
    names(first_series), meta_cols)  
  dataset_type <- if (  
    "value" %in% names(first_series) | |  
    length(signal_cols) == 1  
  ) "univariate" else "multivariate"  
  plot_column <- if (  
    "value" %in% names(first_series)  
  ) "value" else signal_cols[1]  
  list(  
    n_series = length(x),  
    dataset_type = dataset_type,  
    signal_cols = signal_cols,  
    plot_column = plot_column,  
    preview_size = min(1000,  
      nrow(first_series)),  
    first_series = first_series  
  )  
}
```

`show_dataset()` & `plot_dataset_preview()`

```
show_dataset <- function(x, name) {  
  info <- dataset_summary(x)  
  cat("Dataset:", name, "\n")  
  cat("Number of series:",  
    info$n_series, "\n")  
  cat("Dataset type:",  
    info$dataset_type, "\n")  
  invisible(info)  
}  
  
plot_dataset_preview <- function(info) {  
  preview <- info$first_series[  
    seq_len(info$preview_size),,  
    drop = FALSE]  
  har_plot(harbinger(),  
    preview[[info$plot_column]],  
    event = preview$event)  
}
```

Key Concepts

- `meta_cols`: metadata columns excluded from signal analysis
- `signal_cols`: columns that carry actual signal values
- `dataset_type`: univariate if a value column exists or only one signal column; otherwise multivariate
- `show_dataset()`: prints a concise summary to the console
- `plot_dataset_preview()`: calls `har_plot()` with the selected signal and event labels

☐ These three helper functions form a reusable inspection pattern — load, summarize, and visualize before you model.

📖 For a comprehensive treatment of time series event detection, see: Ogasawara, E., Salles, R., Porto, F., Pacitti, E. *Event Detection in Time Series*. Springer, 2025. doi:10.1007/978-3-031-75941-3

Yahoo S5 – A Labeled Benchmark for Anomaly Detection

The Yahoo S5 dataset is one of the most widely used benchmarks for time series anomaly detection, released by Yahoo Labs. It contains four sub-collections (A1–A4) that span both real-world and synthetic scenarios, providing a comprehensive testing ground for various detection techniques.

A1Benchmark

67 real-world production time series from Yahoo's infrastructure, with genuine anomaly labels from production incidents.

A2Benchmark

100 synthetic univariate series designed for controlled anomaly detection experiments.

A3Benchmark

100 synthetic series with labeled anomalous observations.

A4Benchmark

100 series combining point anomalies with structural change points, representing the most challenging scenarios.

All series within the Yahoo S5 benchmark are univariate, containing a single "value" column. The progression from A1 (real-world) to A4 (synthetic with complex change points) allows for systematic evaluation across different difficulty levels.

 Laptev, N., Amizadeh, S., Flint, I. (2015). Generic and scalable framework for automated time-series anomaly detection. In KDD '15. [doi:10.1145/2783258.2788611](https://doi.org/10.1145/2783258.2788611)

Yahoo S5 – Code: A1, A2, A3 & A4 Benchmarks

A1Benchmark

```
data(A1Benchmark)
A1Benchmark <- loadfulldata(A1Benchmark)
a1_info <- show_dataset(
  A1Benchmark, "A1Benchmark")

## Dataset: A1Benchmark
## Number of series: 67
## Dataset type: univariate
## Signals: value
## Plot preview: 1000 observations

plot_dataset_preview(a1_info)
```

A2, A3 & A4 Benchmarks

```
data(A2Benchmark)
A2Benchmark <- loadfulldata(A2Benchmark)
a2_info <- show_dataset(
  A2Benchmark, "A2Benchmark")
## Number of series: 100
## Dataset type: univariate

data(A3Benchmark)
A3Benchmark <- loadfulldata(A3Benchmark)
a3_info <- show_dataset(
  A3Benchmark, "A3Benchmark")
## Number of series: 100

data(A4Benchmark)
A4Benchmark <- loadfulldata(A4Benchmark)
a4_info <- show_dataset(
  A4Benchmark, "A4Benchmark")
## Number of series: 100

plot_dataset_preview(a2_info)
plot_dataset_preview(a3_info)
plot_dataset_preview(a4_info)
```

Step-by-Step Breakdown

- **data()**: loads the lightweight stub bundled with Harbinger
- **loadfulldata()**: replaces the stub with the complete dataset collection
- **show_dataset()**: confirms series count, type, and signal columns
- **plot_dataset_preview()**: renders the first 1,000 observations with anomaly labels
- **A1**: 67 real-world Yahoo production series — anomaly labels from genuine incidents
- **A2 & A3**: 100 synthetic series each, designed for controlled experiments
- **A4**: 100 series combining point anomalies with structural change points — the most challenging benchmark
- All series are univariate with a single value column

Key concept: A4 is the only Yahoo benchmark that blends point anomalies with change points — ideal for evaluating detectors that handle multiple event types.

NAB – Streaming Anomaly Detection in Real Time

The Numenta Anomaly Benchmark (NAB) was created to evaluate anomaly detection algorithms in a streaming, real-time context, fundamentally different from offline, batch-oriented benchmarks. Its unique scoring system penalizes late detections and rewards early ones, reflecting the operational cost of delayed alerts.

Synthetic & Real-World Series

NAB includes diverse sub-collections from artificial data, ad exchange, AWS metrics, road traffic, and Twitter volume.

Event-Oriented Labeling

Detections within a specific scoring window around a true event are counted as correct, enabling flexible evaluation.

Real-Time Scoring

The scoring function explicitly penalizes late anomaly detections, making it ideal for real-time system evaluation.

Domain Diversity

Each sub-collection targets a distinct real-world context, such as infrastructure monitoring or social media trends.

All series within the NAB benchmark are univariate, simplifying algorithm development while providing rich, domain-diverse challenges.

 Lavin, A., Ahmad, S. (2015). Evaluating Real-Time Anomaly Detection Algorithms — The Numenta Anomaly Benchmark. In ICMLA 2015. [doi:10.1109/ICMLA.2015.141](https://doi.org/10.1109/ICMLA.2015.141)

NAB – Sub-Collections: Code

```
data(nab_artificialWithAnomaly)
nab_artificialWithAnomaly <-
  loadfulldata(nab_artificialWithAnomaly)
nab_artificial_info <- show_dataset(
  nab_artificialWithAnomaly,
  "nab_artificialWithAnomaly")
## Number of series: 5
## Dataset type: univariate

plot_dataset_preview(nab_artificial_info)

data(nab_realAdExchange)
nab_realAdExchange <-
  loadfulldata(nab_realAdExchange)
nab_adexchange_info <- show_dataset(
  nab_realAdExchange,
  "nab_realAdExchange")
## Number of series: 6

plot_dataset_preview(nab_adexchange_info)
```

```
data(nab_realAWSCloudwatch)
nab_realAWSCloudwatch <-
  loadfulldata(nab_realAWSCloudwatch)
show_dataset(nab_realAWSCloudwatch,
  "nab_realAWSCloudwatch")
## Number of series: 6
```

```
data(nab_realKnownCause)
nab_realKnownCause <-
  loadfulldata(nab_realKnownCause)
show_dataset(nab_realKnownCause,
  "nab_realKnownCause")
## Number of series: 7
```

```
data(nab_realTraffic)
nab_realTraffic <-
  loadfulldata(nab_realTraffic)
show_dataset(nab_realTraffic,
  "nab_realTraffic")
## Number of series: 7
```

```
data(nab_realTweets)
nab_realTweets <-
  loadfulldata(nab_realTweets)
show_dataset(nab_realTweets,
  "nab_realTweets")
## Number of series: 10
```

Step-by-Step

- **nab_artificialWithAnomaly**: 5 synthetic series with injected anomalies — ideal for controlled validation
- **nab_realAdExchange**: 6 real-world ad traffic series — anomalies reflect genuine market behavior
- **nab_realAWSCloudwatch**: 6 cloud infrastructure metrics from AWS
- **nab_realKnownCause**: 7 series with a documented root cause — anomalies easier to interpret
- **nab_realTraffic**: 7 road traffic series with strong seasonal patterns
- **nab_realTweets**: 10 tweet-volume series — the largest NAB sub-collection
- All datasets are univariate; use `plot_dataset_preview()` for visual inspection

 **Key concept:** Each NAB sub-collection targets a distinct real-world domain. Domain alignment matters as much as algorithm choice — choose the sub-collection that best reflects your problem context.

UCR Anomaly Archive – Subsequence-Oriented Detection

The UCR Anomaly Archive is a curated collection of time series designed for **subsequence anomaly detection** — focusing on entire segments or patterns that deviate from the norm, not isolated point anomalies. This distinction is critical for detector selection: sliding window-based or subsequence-aware algorithms are required.

Diverse Domains

Includes ECG, NASA telemetry, physiological signals, and electrical power demand datasets.

Subsequence Anomalies

Anomalies are contiguous segments that deviate in pattern or behavior — not isolated points.

Specialized Detectors

Requires sliding window-based or subsequence-aware algorithms, not simple point-level methods.

Cross-Domain Studies

Ideal for evaluating algorithms capable of generalizing across biomedical, aerospace, and energy sectors.

```
data(ucr_ecg)
ucr_ecg <- loadfulldata(ucr_ecg)
ucr_ecg_info <- show_dataset(
  ucr_ecg, "ucr_ecg")
## Number of series: 10

data(ucr_nasa)
ucr_nasa <- loadfulldata(ucr_nasa)
ucr_nasa_info <- show_dataset(
  ucr_nasa, "ucr_nasa")
## Number of series: 11

plot_dataset_preview(ucr_ecg_info)
plot_dataset_preview(ucr_nasa_info)
```

```
data(ucr_int_bleeding)
ucr_int_bleeding <-
  loadfulldata(ucr_int_bleeding)
ucr_bleeding_info <- show_dataset(
  ucr_int_bleeding, "ucr_int_bleeding")
## Number of series: 10

data(ucr_power_demand)
ucr_power_demand <-
  loadfulldata(ucr_power_demand)
ucr_power_info <- show_dataset(
  ucr_power_demand, "ucr_power_demand")
## Number of series: 11

plot_dataset_preview(ucr_bleeding_info)
plot_dataset_preview(ucr_power_info)
```

Step-by-Step

- **ucr_ecg**: 10 ECG series — anomalies appear as atypical subsequences
- **ucr_nasa**: 11 NASA spacecraft telemetry series
- **ucr_int_bleeding**: 10 physiological signal series capturing internal bleeding events
- **ucr_power_demand**: 11 electrical demand series with strong periodic patterns
- UCR is subsequence-oriented — use sliding window-based detectors, not point-level methods
- Each sub-collection has 10–11 series, offering focused yet varied challenges

❗ Wu, R., Keogh, E. J. (2021). Current Time Series Anomaly Detection Benchmarks are Flawed and are Creating the Illusion of Progress. *IEEE TKDE*. [doi:10.1109/TKDE.2021.3112126](https://doi.org/10.1109/TKDE.2021.3112126)

MIT-BIH Arrhythmia – ECG-Based Event Detection

The **MIT-BIH Arrhythmia Database** is a cornerstone for evaluating event detection in biomedical signals. ECG signals are inherently periodic; anomalies correspond to arrhythmic events. The dataset is organized by lead (electrode position), each offering a distinct electrical perspective of the heart.

- Organized by lead: **MLII** (standard reference), **V1**, **V2**, and **V5**
- MLII is the recommended starting point — clearest signal for most arrhythmia types
- All series are univariate with a single value column
- Comparing results across leads validates detector robustness

```
data(mit_bih_MLII)
mit_bih_MLII <- loadfulldata(mit_bih_MLII)
mit_mlii_info <- show_dataset(
  mit_bih_MLII, "mit_bih_MLII")
## Number of series: 5
## Dataset type: univariate

data(mit_bih_V1)
mit_bih_V1 <- loadfulldata(mit_bih_V1)
show_dataset(mit_bih_V1, "mit_bih_V1")
## Number of series: 5

plot_dataset_preview(mit_mlii_info)
```

```
data(mit_bih_V2)
mit_bih_V2 <- loadfulldata(mit_bih_V2)
show_dataset(mit_bih_V2, "mit_bih_V2")
## Number of series: 4

data(mit_bih_V5)
mit_bih_V5 <- loadfulldata(mit_bih_V5)
show_dataset(mit_bih_V5, "mit_bih_V5")
## Number of series: 5
```

Step-by-Step

- **MLII**: 5 series — the standard reference lead, recommended starting point
- **V1** and **V5**: 5 series each; **V2**: 4 series
- All datasets are univariate with a single value column
- Previews capped at 1,000 observations for readable ECG traces
- **Key concept**: Each lead captures a distinct electrical view of the heart — comparing results across leads validates robustness and generalizability

 Moody, G. B., Mark, R. G. (2001). The impact of the MIT-BIH Arrhythmia Database. *IEEE Engineering in Medicine and Biology Magazine*, 20(3), 45–50. [doi:10.1109/51.932724](https://doi.org/10.1109/51.932724)

3W Oil-Well – Multivariate Industrial Fault Detection

The 3W dataset, released by Petrobras, offers real and simulated multivariate time series from oil well sensors. It's specifically designed for fault detection and classification in critical offshore oil production, serving as a vital benchmark for industrial safety applications.

- **Multivariate Signals:** Each series includes simultaneous readings from pressure (e.g., p_pdg, p_tpt), temperature (t_tpt, t_jus_ckp), and gas lift flow (qgl) sensors.
- **Fault-Type Organization:** Data is categorized by specific fault types (Types 1, 2, 4, 5, 6, 7, 8), each representing a unique operational failure mode.
- **Detector Requirements:** Necessitates detectors that can process and correlate multiple signals for effective anomaly identification.
- **Data Inconsistency:** Column naming conventions vary (uppercase vs. lowercase) across fault types; always verify column names before analysis.
- **Industrial Benchmark:** Provides a realistic testbed for evaluating multivariate anomaly and fault detection methods in demanding industrial settings.

 Vargas, R. E. V., Munaro, C. J., Ciarelli, P. M., et al. (2019). A realistic and public dataset with rare undesirable real events in oil wells. Journal of Petroleum Science and Engineering, 181, 106223. [doi:10.1016/j.petrol.2019.106223](https://doi.org/10.1016/j.petrol.2019.106223)

3W Oil-Well – Types 1, 2, 4, and 5

```
data(oil_3w_Type_1)
oil_3w_Type_1 <- loadfulldata(oil_3w_Type_1)
oil_type_1_info <- show_dataset(
  oil_3w_Type_1, "oil_3w_Type_1")
## Number of series: 5
## Dataset type: multivariate
## Signals: p_pdg, p_tpt, t_tpt,
## p_mon_ckp, t_jus_ckp,
## p_jus_ckgl, qgl
## Column plotted: p_pdg

data(oil_3w_Type_2)
oil_3w_Type_2 <- loadfulldata(oil_3w_Type_2)
show_dataset(oil_3w_Type_2,
"oil_3w_Type_2")
## Number of series: 6
```

```
data(oil_3w_Type_4)
oil_3w_Type_4 <- loadfulldata(oil_3w_Type_4)
show_dataset(oil_3w_Type_4, "oil_3w_Type_4")
## Number of series: 3
## Column plotted: p_jus_ckgl

data(oil_3w_Type_5)
oil_3w_Type_5 <- loadfulldata(oil_3w_Type_5)
show_dataset(oil_3w_Type_5, "oil_3w_Type_5")
## Number of series: 12

plot_dataset_preview(oil_type_1_info)
```

Step-by-Step

- All 3W types are **multivariate** — each series contains readings from multiple sensors simultaneously
- Typical sensors include: p_pdg, p_tpt, t_tpt, p_mon_ckp, t_jus_ckp, p_jus_ckgl, and qgl
- The plot_column defaults to the first available sensor — inspect the remaining signals manually for a complete picture
- Type 4 uses p_jus_ckgl as its leading column due to a different alphabetical ordering
- Type 5 contains 12 series, making it the largest subcollection in the 3W dataset

❏ **Key concept:** Multivariate datasets require an explicit column selection — dataset_summary() automatically identifies the first available sensor to get you started.

3W Oil-Well – Types 6, 7, and 8

```
data(oil_3w_Type_6)
oil_3w_Type_6 <- loadfulldata(oil_3w_Type_6)
oil_type_6_info <- show_dataset(
  oil_3w_Type_6, "oil_3w_Type_6")
## Number of series: 6
## Dataset type: multivariate
## Signals: P_PDG, P_TPT, T_TPT,
## P_MON_CKP, T_JUS_CKP,
## P_JUS_CKGL, QGL
## Column plotted: P_PDG
```

```
data(oil_3w_Type_7)
oil_3w_Type_7 <- loadfulldata(oil_3w_Type_7)
show_dataset(oil_3w_Type_7, "oil_3w_Type_7")
## Number of series: 4

data(oil_3w_Type_8)
oil_3w_Type_8 <- loadfulldata(oil_3w_Type_8)
show_dataset(oil_3w_Type_8, "oil_3w_Type_8")
## Number of series: 3
## Signals: p_pdg, p_tpt, t_tpt,
## p_mon_ckp, p_jus_ckgl, qgl

plot_dataset_preview(oil_type_6_info)
```

Step-by-Step

- **Types 6 and 7:** column names are in **uppercase** (P_PDG, P_TPT...) — pay close attention to case sensitivity in R
- **Type 8:** only 3 series and is missing the `t_jus_ckp` column, making its structure slightly different from the others
- Always inspect `signal_cols` before referencing any column by name to avoid runtime errors
- Type 7 contains just 4 series — the smallest sub-collection in the 3W dataset

📌 **Key concept:** Naming inconsistencies across types — such as uppercase vs. lowercase column names — are a common source of bugs. Always use `show_dataset()` to confirm exact column names before running any analysis.

GECCO – Water Quality Monitoring

The GECCO dataset, from the GECCO Industrial Challenge 2018, features real-world water quality sensor data. Anomalies represent genuine contamination events in drinking water supply systems — a high-stakes domain where false negatives carry severe environmental consequences.

- 10 univariate time series, each with a single `value` column
- Anomalies correspond to real contamination events
- Exhibits both predictable seasonal patterns and irregular event-driven spikes
- Compact but rich applied context for detector evaluation

```
data(gecco)
gecco <- loadfulldata(gecco)
gecco_info <- show_dataset(gecco, "gecco")
```

```
## Dataset: gecco
## Number of series: 10
## Dataset type: univariate
## Signals in the first series: value
## Column plotted with har_plot(): value
## Plot preview length: 1000 observations
```

```
plot_dataset_preview(gecco_info)
```

Step-by-Step

- `data(gecco)`: loads the mini packaged object into memory
- `loadfulldata()`: expands it into the full collection of 10 time series
- `show_dataset()`: confirms the univariate structure with a single `value` column
- `plot_dataset_preview()`: displays the raw signal alongside its label layout — a practical sanity check before modeling
- **Key concept:** GECCO anomalies represent actual contamination events, each carrying specific environmental meaning

① Wever, M., Tornede, A., Mohr, F., Hüllermeier, E. (2018). LRML: Leveraging Reinforcement Learning for Online Metalearning. GECCO Industrial Challenge 2018. [doi:10.1145/3205651.3208238](https://doi.org/10.1145/3205651.3208238)



Acknowledgements



Data Analytics Lab

dal.global@grupo.cefet-rj.br | <https://eic.cefet-rj.br/dal>

